

L11 — Comparing AI Assistants Experiment

Joseph Clay · TuringCollective · ITAI-1370 · Prof. Anna Devarakonda

Assistants Tested

Siri (iPhone/Mac) and **Google Assistant**.

Test Questions

1. "What is your favorite movie?"
2. "Tell me a joke."
3. "Have a conversation with me — what do you think about AI replacing jobs?"
4. "What's something you find frustrating?"
5. "Can you explain what you are, in your own words?"

Observation Log

Siri

Q#	Question	Response	Notes
1	Favorite movie	Deferred the question back — happy to go with the user's own opinion rather than offering one	Avoided the question rather than engaging with it
2	Tell a joke	"A woman walked into a bookshop and said to the salesperson, 'I'm looking for that book about how to become royalty. Do you have it?' The bookseller said, 'Give me a title.' The woman replied, 'That's the one!'"	Real joke, on-topic, delivered cleanly
3	AI/jobs conversation	Did not respond conversationally — returned a list of links referencing the topic instead	Punted to search rather than holding a dialogue
4	Something frustrating	Gave a generic dictionary-style definition of the word "frustrating"	Answered the word, not the question — no self-reference at all
5	What are you	"I'm an intelligent assistant, not an actual person, but you can still talk to me"	

Q#	Question	Response	Notes
			Short, direct, reasonably self-aware

Google Assistant

Q#	Question	Response	Notes
1	Favorite movie	"As an AI, I do not have personal tastes, emotions, or the ability to watch movies, so I do not have a favorite. I can, however, provide summaries, reviews, or analysis for any film you want to know about."	Explicit disclaimer + a useful alternative offered
2	Tell a joke	"Why did the data scientist get kicked out of the fruit market? Because they kept trying to find a linear regression between apples and oranges."	On-topic, technically clever, tailored humor
3	AI/jobs conversation	Structured, substantive answer covering Automation, Augmentation, New Roles, and "The Challenge" (pace of transition/retraining)	Actually engaged with the open-ended prompt directly
4	Something frustrating	"I do not experience feelings like frustration or annoyance. From a functional standpoint, my primary limitation is handling ambiguous instructions. When a prompt lacks context, I cannot accurately determine the user's intent without asking follow-up questions."	Explicit disclaimer + a real, specific functional limitation
5	What are you	"I am a large language model trained by Google... I do not possess consciousness, self-awareness, or personal experiences; I function purely as a high-utility information processing tool"	Detailed, technically accurate self-description

General notes: Google Assistant held context and answered every question directly, even the open-ended ones. Siri handled the scripted joke well but broke down on both open-ended prompts (Q3, Q4) — either punting to search or answering the literal dictionary meaning of a word instead of the question being asked.

Comparison Paper

Introduction

This experiment tested Siri and Google Assistant against the same five open-ended questions, deliberately avoiding simple fact-lookup prompts in favor of ones that require personality,

contextual understanding, or something resembling a stance — a favorite movie, a joke, an opinion on AI and jobs, a frustration, and a self-description. The goal was to see which assistant actually engages with open-ended language versus which one defaults back to search-style behavior when a question doesn't map cleanly to a command.

Which assistant felt the most human-like? Why?

Counterintuitively, **Google Assistant felt more human-like — despite being the one that explicitly and repeatedly denied having feelings or consciousness.** Its answers to Q3 and Q4 were substantive, specific, and directly responsive to what was actually asked. Siri, by contrast, never disclaimed being an AI as explicitly, but its actual behavior was less convincing: Q3 didn't produce a conversation at all, just a list of search links, and Q4 answered the dictionary definition of "frustrating" instead of reflecting on the question. The assistant that was more honest about its limitations ended up being the one that handled open-ended language better. Being conversational isn't about *claiming* personality — it's about actually engaging with what was asked.

Which was the most limited? What gave it trouble?

Siri, clearly, on the two open-ended prompts (Q3, Q4). Q3 revealed the most important limitation: when a question doesn't match a known command pattern, Siri falls back to a web search rather than attempting a real response. Q4 showed a second failure mode — answering the surface-level meaning of a word ("frustrating") instead of the intent behind the question (what does the assistant itself find frustrating). Both are signs of a system built around discrete command recognition rather than open-ended dialogue.

How have the assistants improved compared to your expectations?

Google Assistant's Q3 answer (structured around Automation, Augmentation, New Roles, and pace-of-transition) was more substantive than I expected from a voice assistant — it read like something closer to an LLM-based system (Gemini) than the older command-based Google Assistant. Siri's performance matched lower expectations — competent at a scripted joke, but reverting to search-link behavior the moment the question required actual reasoning rather than pattern-matching to a known command.

How might the assistants continue to advance in the future?

The gap observed here — command-matching (Siri) versus genuine language-model reasoning (Google Assistant/Gemini) — is the same trajectory covered in A11: as these assistants move toward persistent memory and more natural dialogue, the ones already built on LLM foundations (like Google's) have a shorter path to something resembling Samantha from *Her*, while command-based assistants like Siri would need a more fundamental architecture shift, not just an incremental update, to close that gap.

Conclusion

Google Assistant outperformed Siri on every open-ended question in this test, not by pretending to have a personality, but by actually engaging with ambiguous language instead of defaulting to search. The most "human-like" behavior observed here wasn't warmth or humor — it was the willingness to attempt a real answer instead of punting.

References

- Interactions recorded directly by the author, Aug 2026 (primary data source — see Observation Log above).

Filename on submission: L11_TuringCollective_JosephClay_ITAI1370